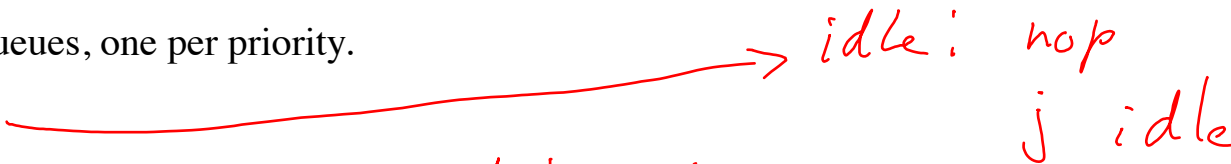


Phase 1: Notes and Hints

- Getting started
 - The USLOSS source is available at `/home/cs452/spring18/usloss/usloss3.6.tgz`
 - Compiling USLOSS: See Section 2.1 of the USLOSS User's Manual
 - Some files to help you get started can be found in `/home/cs452/spring18/phase1`. Copy these to your own phase directory.
 - Your phase1 code will be compiled into a library named: **libphase1.a** in your directory.
 - To execute a test case, you link the `.o` file of the test case with the **libusloss<version>.a** and **libphase1.a** libraries.
 - Typing `'make'` will create the **libphase1.a** library in your directory.
 - Typing `'make test00'`, for example, will create an executable test case named `'test00'` in your directory.

- Header files for Phase 1:
 - After you have compiled USLOSS:
 - **phase1.h**: contains function prototypes & constants to be used in this phase.
 - **~/include/usloss.h**: contains function prototypes for USLOSS library functions, many useful constants.
 - **kernel.h**: your data structures and constants for phase1.
- Mode and interrupts
 - All functions in this phase have to be executed in kernel mode.
 - Test for kernel mode must be done for each phase 1 function, since processes must be in kernel mode to call these.
 - See discussion of Process Status Register (Section 3.4) and functions to access this register (Section 7) in the USLOSS manual.
 - Some useful constants are defined in **usloss.h**.
 - Enabling and disabling interrupts:
 - Manipulate the appropriate bits in the PSR register.
 - Interrupts can be turned on and off only in kernel mode
 - When should interrupts be disabled? (A key point in writing correct phase 1 functions.)

- **startup()** function

- Called by USLOSS. Note: **main()** is inside the USLOSS library, not in your phase 1 code.
- Initialize your kernel data structures, in particular, the process table.
 - The **struct procStruct** in the provided **kernel.h** is not complete. Add/change fields as needed!
- Ready List(s): for processes waiting to run. Choose one of:
 - Single queue arranged by priorities; or
 - Multiple queues, one per priority.
- Sentinel process: 
 - Call **fork1()** to create the sentinel process. *don't call dispatcher*
 - Sentinel runs when there are no more runnable processes.
 - It is the only process with priority **6**; its status is always READY.
 - The provided **sentinel()** function in **skeleton.c** is complete.
 - You will need to provide the **checkDeadlock()** function that **sentinel()** calls.
- Call **fork1()** to create a process that will execute **start1()**.
 - **start1()** will be inside every phase 1 test case.

fork1() function:

- Test for kernel mode.
- Initialize process table entry, which will include:
 - Check for valid stack size and priorities; there is a minimum stack size.
 - Use **malloc()** (or **posix_memalign()**) to create the stack. Save stack information in process table so stack can be freed in **quit()**.
 - This will be the only call to **malloc** (or **posix_memalign**) in phases 1, 2, 3, and 4!!
 - You will get to dynamically allocate memory again in phase 5.
- Store function pointer and argument value:
 - **fork1()** does not execute the function.
 - **launch()** does.
- Initialize context using the USLOSS **USLOSS_ContextInit()** function.
 - Use **launch()** as the function pointer passed to **USLOSS_ContextInit()**.
 - Necessary: must enable interrupts before starting function, and correctly handle return from function.
 - Unix analogy: What happens if **main** calls **return** instead of **exit**?
- Call the **dispatcher()** — let the **dispatcher()** decide whether the parent or the child runs next!
- Enable interrupts (for the parent) and return the pid of the child process.

join() function: There are three cases:

- The process has no children. (What happens?)
- Child(ren) **quit()** before the **join()** occurred.
 - Return the pid and quit status of one quit child and finish the clean up of that child's process table entry.
 - Report on quit children in the order the children **quit()**.
- No (unjoined) child has **quit()** ... must wait.
 - How?
 - After wait is over: return the pid and quit status of the child that quit.
 - Where does the parent find these?
- The child status that **join** returns is the argument that the child passed to **quit()**.
- HINT: Pick one of the **3** cases above and work on that -- ignore the other 2.
 - HINT: look at test00.c. Which case does it fall into? Implement that case first!

quit() function

- Error if a process with active children calls **quit()**. Halt USLOSS with appropriate error message.
- Cleanup the process table entry (but not entirely, see **join()**)
 - Two cases:
 - Parent has already done a **join()**, or
 - Parent has not (yet) done a **join()**.
- Unblock processes that **zap**'d this process (see below).
- May have children who have **quit()**, but for whom a **join()** was not (and now will not) be done.
 - This is not an error.

dispatcher() function: Decides which process gets to run.

- When is the dispatcher called?
- Checks if the current process can continue running: Has it been time-sliced? Has it been blocked? Is it still the highest priority among **READY** processes?
- Select a new process and perform a context switch in order to get it running.
- Choose according to scheduling policy: see Section 3.2 of Phase 1 handout.

Clock interrupt handler (see Section 3.2 of USLOSS manual).

- Defer working on this until your **fork1**, **join**, **quit**, and **dispatcher** functions are all working.
- Write an **clockHandler** function. Its prototype will be:

```
void clockHandler(int dev, void *arg)
```

- In the code for **clockHandler**, you will ignore the **arg** parameter; it is not used by the clock.
- Interrupt vector is defined by USLOSS as an array of pointers to **void** functions with 2 integer arguments:

```
extern void (*USLOSS_IntVec[NUM_INTS]) (int dev, void *arg); /* from usloss.h */
```

- Initialize the appropriate slot to point to your clock handler function.

```
USLOSS_IntVec[USLOSS_CLOCK_DEV] = clockHandler;
```

- The other slots can be ignored (they will show up again in later phases).
- **clockHandler()** function.
 - Checks if the current process has exceeded its time slice. Calls **dispatcher()** if necessary.
 - Time slice is 80 ms (milliseconds).
 - The **USLOSS_Clock()** function returns time in microseconds (= 1,000 ms); thus, time slice is 80,000 μ s.

zap() and **isZapped()** functions.

- The zap'ing process blocks until the zap'd process quits.
- Store information in the process table of both the zap'ing and zap'd processes.
- Zap'ing process blocks itself. (How?)
- The zap'd process, when it quits, needs to know about the zap'ing process.
- Note: there can be more than one zap'ing process! During a **quit()**, unblock all zap'ing processes.


```

PREFIX = ${HOME}

TARGET = libphase1.a
ASSIGNMENT = 452phase1
CC = gcc
AR = ar

COBJS = phase1.o
CSRCS = ${COBJS:.o=.c}

HDRS = kernel.h phase1.h

INCLUDE = ${PREFIX}/include

CFLAGS = -Wall -g -I${INCLUDE} -I. -std=gnu99

UNAME := $(shell uname -s)

ifeq ($(UNAME), Darwin)
    CFLAGS += -D_XOPEN_SOURCE           # use for Mac, NOT for Linux!!
endif

LDFLAGS = -L. -L${PREFIX}/lib

TESTDIR = testcases
TESTS = test00 test01 test02 test03 test04 test05 test06 test07 test08

LIBS = -lphase1 -lusloss3.6

```

```
$(TARGET): $(COBJS)
    $(AR) -r $@ $(COBJS)

$(TESTS): $(TARGET) p1.o
    $(CC) $(CFLAGS) -c $(TESTDIR)/$.c
    $(CC) $(LDFLAGS) -o $@ $.o $(LIBS) p1.o

clean:
    rm -f $(COBJS) $(TARGET) p1.o test?.o test? test?.txt core term*.out

phases1.o: kernel.h

submit: $(CSRCS) $(HDRS) Makefile
    tar cvzf phases1.tgz $(CSRCS) $(HDRS) Makefile
```

Selected parts of `/home/cs452/spring18/phase1/skeleton.c`

```
void startup(int argc, char *argv[])
{
    int i;          /* loop index */
    int result; /* value returned by call to fork1() */

    /* initialize the process table */

    /* Initialize the Ready list, etc. */
    if (DEBUG && debugflag)
        USLOSS_Console("startup(): initializing the Ready & Blocked lists\n");
    ReadyList = NULL;

    /* Initialize the clock interrupt handler */

    /* startup a sentinel process */
    if (DEBUG && debugflag)
        USLOSS_Console("startup(): calling fork1() for sentinel\n");
    result = fork1("sentinel", sentinel, NULL, USLOSS_MIN_STACK, SENTINELPRIORITY);
    if (result < 0) {
        if (DEBUG && debugflag)
            USLOSS_Console("startup(): fork1 of sentinel returned error, halting...\n");
        USLOSS_Halt(1);
    }
}
```

```

/* start the test process */
if (DEBUG && debugflag)
    USLOSS_Console("startup(): calling fork1() for start1\n");
result = fork1("start1", start1, NULL, 2 * USLOSS_MIN_STACK, 1);
if (result < 0) {
    USLOSS_Console("startup(): fork1 for start1 returned an error, halting...\n");
    USLOSS_Halt(1);
}

USLOSS_Console("startup(): Should not see this message! ");
USLOSS_Console("Returned from fork1 call that created start1\n");

return;
} /* startup */

int sentinel (char *dummy)
{
    if (DEBUG && debugflag)
        USLOSS_Console("sentinel(): called\n");
    while (1)
    {
        checkDeadlock();
        USLOSS_WaitInt();
    }
} /* sentinel */

```

```

int fork1(char *name, int (*startFunc)(char *), char *arg, int stacksize, int priority)
{
    int procSlot;

    if (DEBUG && debugflag)
        USLOSS_Console("fork1(): creating process %s\n", name);

    // test if in kernel mode; halt if in user mode

    // Return if stack size is too small

    // Is there room in the process table? What is the next PID?

    // fill-in entry in process table
    if ( strlen(name) >= (MAXNAME - 1) ) {
        USLOSS_Console("fork1(): Process name is too long.  Halting...\n");
        USLOSS_Halt(1);
    }
    strcpy(ProcTable[procSlot].name, name);
    ProcTable[procSlot].startFunc = startFunc;
    if ( arg == NULL )
        ProcTable[procSlot].start_arg[0] = '\0';
    else if ( strlen(arg) >= (MAXARG - 1) ) {
        USLOSS_Console("fork1(): argument too long.  Halting...\n");
        USLOSS_Halt(1);
    }
    else
        strcpy(ProcTable[procSlot].startArg, arg);
}

```

```

// Initialize context for this process, but use launch function pointer for
// the initial value of the process's program counter (PC)

USLOSS_ContextInit(&(ProcTable[procSlot].state),
                  ProcTable[procSlot].stack, ProcTable[procSlot].stackSize,
                  NULL, launch);

/* for future phase(s) */
pl_fork(ProcTable[procSlot].pid);

// More stuff to do here...

} /* fork1 */

```

```

void launch()
{
    int result;

    if (DEBUG && debugflag)
        USLOSS_Console("launch(): started\n");

    // Enable interrupts

    // Call the function passed to fork1, and capture its return value
    result = Current->startFunc(Current->startArg);

    if (DEBUG && debugflag)
        USLOSS_Console("Process %d returned to launch\n", Current->pid);

    quit(result);
} /* launch */

```